

# Beyond Keywords: AI-powered Text Search with pgvector for PostgreSQL



**Belma Canik**

(She/Her)

Senior Database Specialist TAM



# Agenda

GenAI


Vector concept

PostgreSQL as vector store with pgvector

Demo lab: AI powered similarity search using pgvector


Questions

# Generative AI




## Artificial intelligence (AI)

Any technique that allows computers to mimic human intelligence using logic, if-then statements, and machine learning



## Machine learning (ML)

A subset of AI that uses machines to search for patterns in data to build logic models automatically



## Deep learning (DL)

A subset of ML composed of deeply multi-layered neural networks that perform tasks like speech and image recognition



## Generative AI

Powered by large models that are pre-trained on vast corpora of data and commonly referred to as foundation models (FMs)

# Generative AI is powered by foundation models (FMs)

Pre-trained on vast amounts of unstructured data

---


Contain large number of parameters that make them capable of learning complex concepts

---

Can be applied in a wide range of contexts

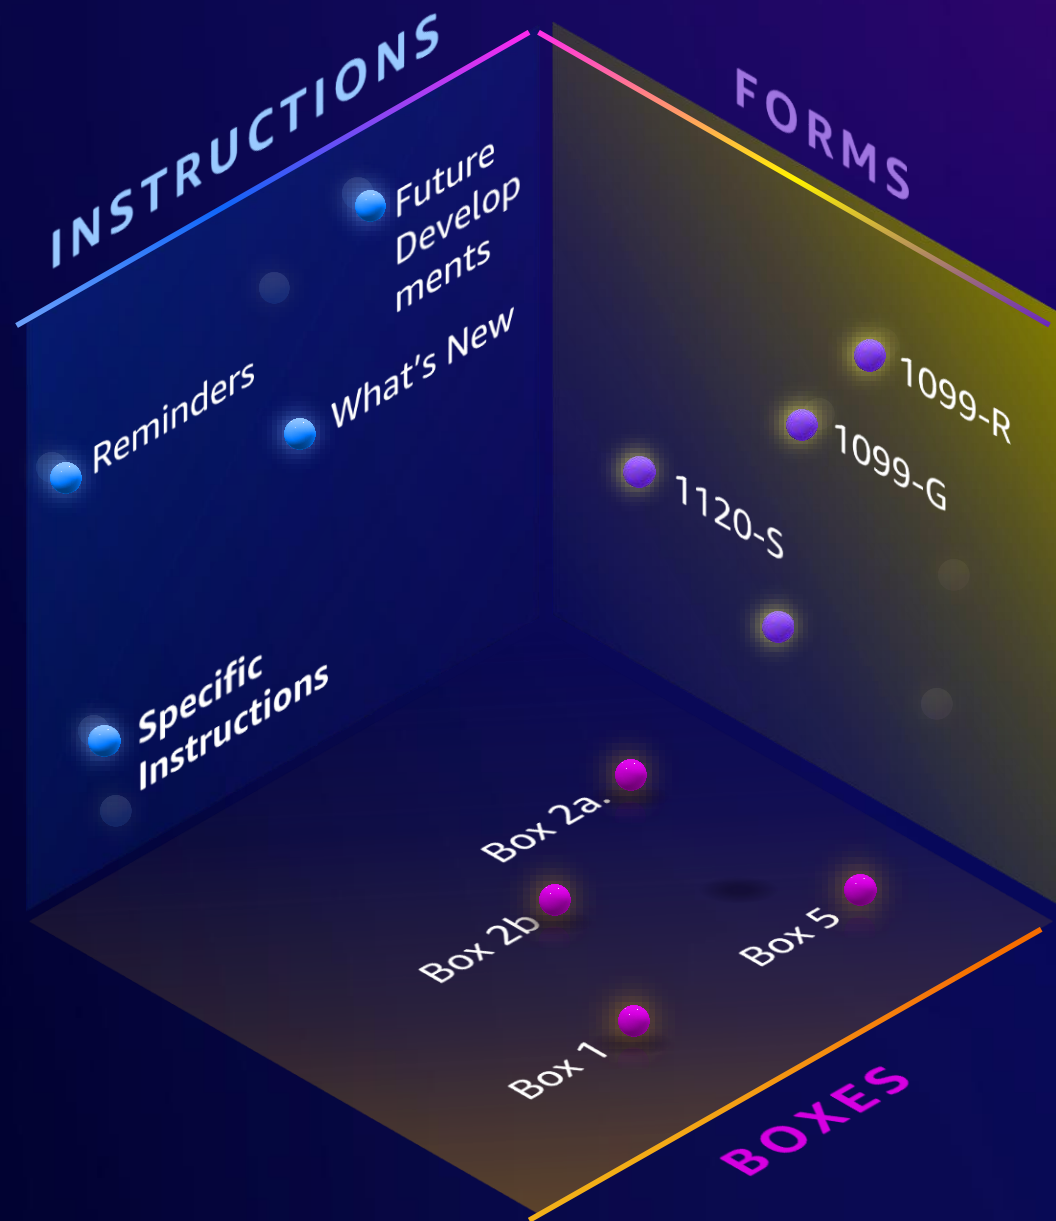
---

Customize FMs using your data for domain-specific tasks




# Vector space

VECTORS ARE ORGANIZED  
IN CLUSTERS



# What is a **vector embedding** ?

- A **numerical representation** of words or sentences, used in NLP
- NLP models can easily perform tasks such as **querying, classification, and applying machine learning algorithms** on textual data




# What are vector embeddings?



Unstructured data has to be vectorized into vectors to be used in generative AI applications



# An example of Generated embedding in vector space







# PostgreSQL as a vector store

# Why use PostgreSQL for vector searches?



Existing client libraries work without modification



Convenient to co-locate app and AI/ML data in same database



PostgreSQL acts as persistent transactional store while working with other vector search systems

**Note:** Postgres, PostgreSQL, and the Slonik Logo are trademarks or registered trademarks of the PostgreSQL Community Association of Canada, and used with their permission

# Native vector support and challenges

## ARRAY data type

Multiple data types (int4, int8, float4, float8)  
“Unlimited” dimensions  
No native distance operations  
Can add using Trusted Language Extensions + PL/Rust  
No native indexing

## Cube data type

- float8 values
- Euclidean, Manhattan, Chebyshev distances
- K-NN GiST index – exact nearest neighbor search
- Limited to 100 dimensions

# What is pgvector?

- An open-source extension
- support for **storage, indexing, searching, metadata** with choice of **distance**

**vector** data types

Co-locate with embeddings

Exact nearest neighbor  
Approximate nearest neighbor (ANN)

Supports **IVFFlat/HNSW** indexing

Distance operators (**<->**, **<=>**, **<#>**, **<+>**, **<~>**, **<%>**)


[github.com/pgvector/pgvector](https://github.com/pgvector/pgvector)

Note: **<+>**, **<~>**, and **<%>** operators available only from pgvector version 0.7.0

Aurora PostgreSQL 15.3, 14.8, 13.11, 12.15 and higher  
Amazon RDS PostgreSQL 15.2 and higher



# Pgvector Popularity




# Approximate nearest neighbor (ANN)

- Find similar vectors without searching all of them

- Faster than exact nearest neighbor


- “Recall” – % of expected results



# pgvector: offers distance operations

## Euclidean (L2) (vector\_l2\_ops)

Useful for counts / measurements  
Recommendation Systems




$$d(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{i=1}^n (q_i - p_i)^2}$$

<->

## Cosine Similarity (vector\_cosine\_ops)

Useful for semantic search and  
document classification




$$\text{sim}(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \cdot \|\mathbf{b}\|}$$

<=>

## Dot Product (vector\_ip\_ops)

Useful for collaborative filtering



$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos \alpha$$

<#>



# pgvector: new operations from version 0.7.0

Manhattan Distance (L1) aka Taxicab  
(vector\_l1\_ops)

HNSW indexing

Useful for city navigation /  
speech recognition and image  
processing




$$Manhattan = \sum_{i=1}^n |x_i - y_i|$$

<+>

Hamming (binary vectors)  
(bit\_hamming\_ops)

Useful for error correction / dissimilarity  
between two binary vectors or strings




sum(vectOr1 != vector2)

<~>

Jaccard Distance(binary vectors)  
(bit\_jaccard\_ops)

Image recognition with labelled data /  
Compare text patterns in documents  
based on the overlap of words.



$$Jaccard = \frac{\text{Intersection (A, B)}}{\text{Union (A, B)}}$$

<%>



# Inverted File with Flat Compression (IVFFlat) Index

# Building an IVFFlat index




# Building an IVFFlat index: Find centers



# Building an IVFFlat index: Assign lists




# Querying an IVFFlat index



```
SELECT id FROM products ORDER BY $1 <-> embedding LIMIT 3
```

# Querying an IVFFlat index



```
SELECT id FROM products ORDER BY $1 <-> embedding LIMIT 3
```




# IVFFlat query parameters

- `ivfflat.probes` (1 by default)
  - Number of lists to search during a query
  - More lists leads to more relevant results
  - More lists requires more time
  - $\text{probes} = \sqrt{\text{lists}}$

# Best practices for building IVFFlat indexes

- Choose value of lists to maximize recall but minimize effort of search
  - < 1MM vectors:  $\# \text{ vectors(rows)} / 1000$
  - > 1MM vectors:  $\sqrt{\# \text{ vectors(rows)}}$
- May be necessary to rebuild when adding/modifying vectors in index
- Use parallelism to accelerate build times
- Increase maintenance\_work\_mem for faster index creation
- Faster build times, less memory but lower query performance

# Using parallelism to accelerate IVFFlat builds



# Performance strategies for IVFFlat queries


- Increasing `ivfflat.probes` increases recall, decreases performance
- Lowering `random_page_cost` on a per-query basis can induce index usage
- Set `shared_buffers` to a value that keeps data (table) in memory
- Increase `work_mem` on a per-query basis



# Hierarchical Navigable Small World (HNSW) Index

# HNSW (Hierarchical Navigable Small Worlds)

- a multi-layered graph-based indexing
- the skip lists and navigable small world algorithms



Source: <https://tembo.io/blog/vector-indexes-in-pgvector/#hns>


<https://jkatz05.com/post/postgres/pgvector-hnsw-performance/>

# HNSW index building parameters


- **m**
- Maximum number of bidirectional links between indexed vectors
  - Default: 16
- **ef\_construction**
- Number of vectors to maintain in “nearest neighbor” list
- Default: 64



# Choosing m and ef\_construction (serial)



# Choosing m and ef\_construction (parallel)



# Best practices for HNSW indexes

## Building HNSW indexes

- Default values (`m=16, ef_construction=64`) usually work
- (pgvector 0.5.1) Start with empty index and use concurrent writes to accelerate builds

## Performance strategies

- Index building has biggest impact on performance/recall
- Increasing `hnsw.ef_search( default 40)` increases recall, decreases performance

# Which index do I choose?

If you care more about index size, then choose IVFFlat

If you care more about index build time, then select IVFFlat

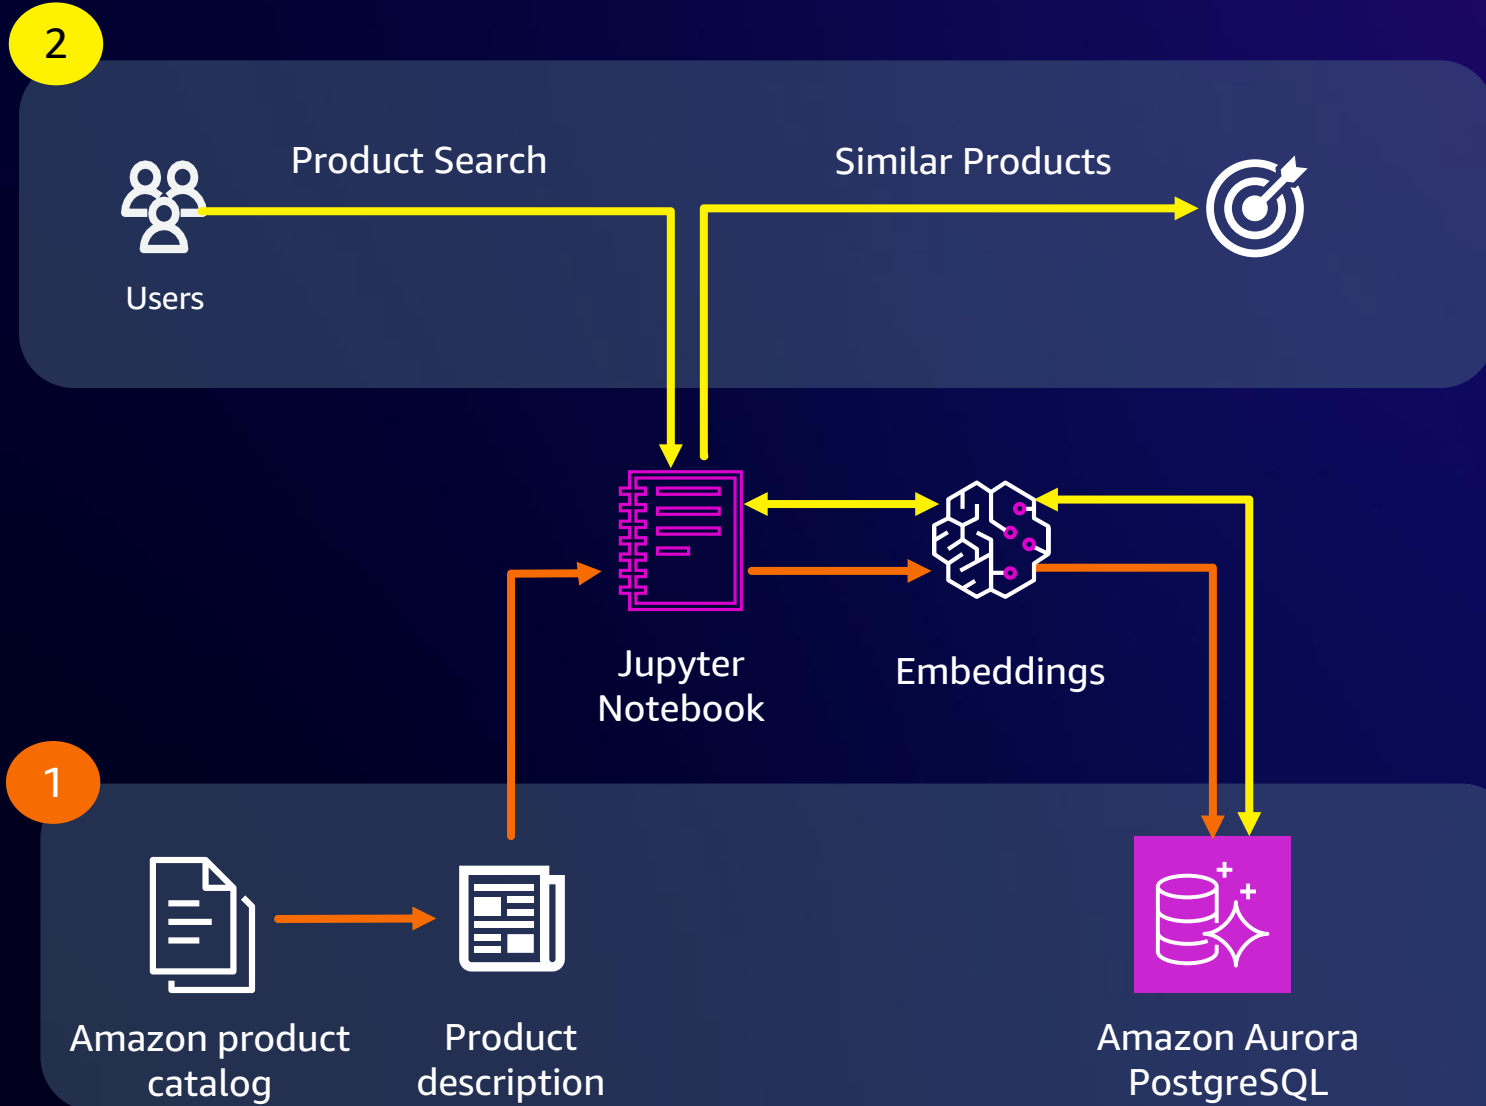
If you care more about high performance/recall, then choose HNSW

If you expect vectors to be added or modified, then select HNSW



# Demo lab: AI powered similarity search using pgvector

# How vector embeddings are used



# Questions



# Thank You

**Belma Canik**

<https://www.linkedin.com/in/belma-canik-38b66648/>

