



Getting started with pgwatch
Prague PostgreSQL Meetup:
February Edition

Senior Consultant/Developer

Pavlo Golub

MAIL

pavlo.golub@cybertec.at

Social

pashagolub.github.io

WEB

www.cybertec-postgresql.com



DATABASE - PRODUCTS





AGENDA

- Different levels of database monitoring
- PostgreSQL monitoring approaches
- Features of pgwatch
- Use cases

Getting started with pgwatch



Different levels of database
monitoring

Why to monitor?

- Failure / Downtime detection
- Slowness / Performance analysis
- Proactive predictions
- Maybe wasting money?



Different levels of database monitoring

- High level service availability
- System monitoring
- PostgreSQL land





High level service availability

- Try to periodically connect/query from an outside system
- DIY - e.g. a simple Cron script
- SaaS - lots of service providers

Who will guard the guards?

- You'll probably want two services for more critical stuff...

System monitoring

Too many tools, no real standards. Just make sure to understand what you're measuring!

- Do you know what does the CPU load number actually mean?
 - Is it a good metric?
- What's the difference between VIRT, RES, SHR memory values for a process?



PostgreSQL land

- Log analysis
- Statistics Collector
- Extensions



Log analysis

- “Just in case” storing of logs for possible ad hoc needs
 - Moving logs to a central place makes sense
 - rsync + Cron
- Active parsing
 - grep + Cron
 - DIY (file_fdw, Graylog, ELK, ...)
 - pgBadger (JSON format)
 - Grafana Loki
 - Some cloud service (Loggly, Splunk, ...)

Logging configuration

- Some settings to note
 - log_destination (CSV format recommended)
 - log_statement = 'none' (default)
 - log_min_duration_statement / log_duration
 - log_min_messages / log_min_error_statement

```
postgres=# SELECT count(*) FROM pg_settings  
WHERE category LIKE 'Reporting and Logging%';
```

count

Statistics Collector

- Not all `track_*` parameters enabled by default
- Dynamic views
 - `pg_stat_activity`, `pg_stat_(replication|wal_receiver)`
 - `pg_locks`, `pg_stat_ssl`, `pg_stat_progress_*`
- Cumulative views
 - Most `pg_stat_*` views
 - Long uptimes cause “lag” for problem detection
- Selective stats reset possible

Extensions

- Most notably **pg_stat_statements** (“top statements”)
- **pg_stat_monitor** (pg_stat_statements on steroids)
- **pgstattuple** (bloat)
- **pg_buffercache** (what’s in the shared buffers)
- **auto_explain** (e.g. to analyze “jumping runtimes”)
- **pg_qualstats** (WHERE predicate stats)
- **pg_stat_kcache** (to sample O/S metrics)
- ...

PostgreSQL Monitoring Tools

FIND MORE



<https://wiki.postgresql.org/wiki/Monitoring>

No shortage of tools!

Getting started with pgwatch



Features

Thanks!



Kaarel Moppel

The author of pgwatch v2



Main principles - why another tool?

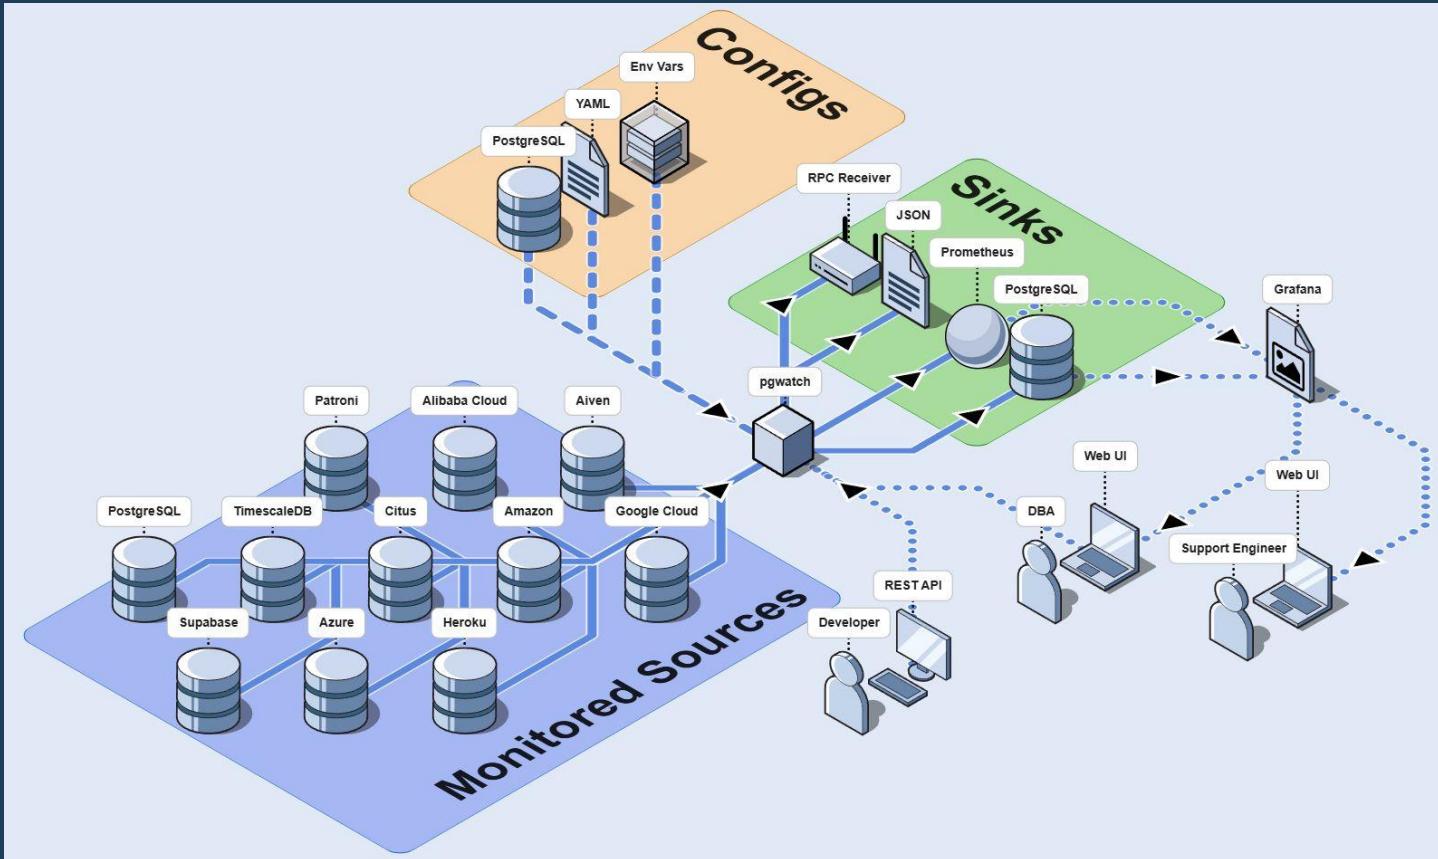
- 1-minute setup
 - One binary!
- User changeable visuals / dashboarding
- Non-invasive
 - No extensions or superuser needed for base functionality
- Easy extensibility, do minimal work needed
 - SQL metrics
- Simple alerting via Grafana possible



Architecture components

- Metrics gathering daemon (Golang)
- Configuration
 - PostgreSQL compatible database
 - YAML files
 - ENV vars
- Measurements sinks
 - PostgreSQL (with TimescaleDB optionally)
 - Prometheus
 - Local file
 - RPC
- Optional simple Web UI for administration
- Grafana for intuitive point-and-click dashboarding

Architecture components





Features

- “Ready to go”
 - Default metrics cover all pg_stat* views
 - Pre-configured dashboards for almost all metrics
- Supports Postgres 11+ out of the box
- Configurable security - passwords, SSL
- Reuse of existing Postgres, Grafana, Prometheus installations possible
- Kubernetes/OpenStack ready

Features

- Per DB setup with optional auto-discovery of all DBs of a cluster
- Change detection for table/index/sproc/configuration events
- AWS RDS CloudWatch metrics support
- PgBouncer & Pgpool2 metrics support
- Patroni support
- Very low resource requirements
- A Ping mode to test connectivity to monitored databases
- Extensible
 - Custom metrics via SQL, i.e. usable also for business layer!
 - Built-in log parsing and OS level metrics
 - Grafana has lot of plugins as well

Alerting

- Quite easy with Grafana, “point-and-click”
- Alertmanager if Prometheus envolved
- Use **pg_timetable** scheduler to analyze and send alerts

Among other things

- TimescaleDB as a metric storage is available out of the box
 - compression is on
 - 45M rows of real world stat_statements data (6db * 3mo)
 - current storage: 70 GB → 8 GB compressed (ratio 8.9x)
 - upcoming feat: 11 GB → 763 MB (ratio 14.5x)

Among other things

- Parallel measurement storages

```
src(master) ~ ?1 1.22.0 ➤ ./pgwatch3 --config=sources/sample.sources.yaml  
--sink=jsonfile://metrics.json --sink=postgresql://pgwatch3@localhost/pgwatch3_metrics --  
sink=prometheus://:9127  
2024-04-11 13:00:03.544 [INFO] [sink:jsonfile://metrics.json] measurements sink activated  
2024-04-11 13:00:03.683 [INFO] [sink:postgresql://pgwatch3@localhost/pgwatch3_metrics] initialising the measurement database ...  
2024-04-11 13:00:03.690 [INFO] [sink:postgresql://pgwatch3@localhost/pgwatch3_metrics] measurements sink activated  
2024-04-11 13:00:03.690 [INFO] [sink:prometheus://:9127] measurements sink activated  
2024-04-11 13:00:03.691 [INFO] [metrics:73] [sources:1] host info refreshed  
2024-04-11 13:00:03.867 [INFO] Connect OK. [test1] is on version 16.1 (in recovery: false)  
2024-04-11 13:00:03.867 [INFO] Trying to create helper functions if missing for "test1" ...  
2024-04-11 13:00:03.952 [INFO] [sink:postgres] [rows:1] [db:pgwatch3_metrics] [elapsed:5.83  
72ms] measurements written
```

Among other things

- Docker images
 - are multiplatform and are slimmer
 - no more monstrous all-in-one image
 - docker compose is shipped instead

```
docker images -a
```

REPOSITORY	TAG	IMAGE ID	SIZE
cybertec/pgwatch2-postgres	latest	7a889edafe50	1.17GB
cybertecpostgresql/pgwatch-demo	latest	5e3b24fb5a6a	753MB
cybertecpostgresql/pgwatch	latest	2cbd8fc36e28	44.1MB

Getting started with pgwatch



Use cases

Use case 1: Live demo

1. docker run -d --name pw3 \
-p 5432:5432 -p 3000:3000 -p 8080:8080 \
-e PW_TESTDB=true \
cybertecpostgresql/pgwatch-demo
2. Wait some seconds and open browser at localhost:8080
3. Insert your DB connection strings
4. Start viewing/editing dashboards in 5 min...

or check online demo at **<https://demo.pgwatch.com>**



Use case 2: Find slow queries

1. pgwatch is running with “stat_statements” metric enabled
2. “pg_stat_statements” extension enabled in monitored database
3. Data collected for some time after an hour of regular workload or after peak hours
4. Examine “Stat Statements Top” dashboards



Use case 3: Detect performance spikes

1. pgwatch is running with system and PostgreSQL resource metrics enabled.
2. High-load event occurs (e.g., sudden CPU, memory, or I/O spike).
3. Examine dashboards:
 - System Stats → CPU, memory, and disk I/O usage
 - DB Overview → General database activity
 - Sessions Overview → Detect connection spikes
 - Stat Statements Top / Visual → Identify queries causing load
 - Checkpointer / Bgwriter / Block IO Stats → Spot I/O-heavy operations
4. Correlate findings to pinpoint the root cause and take action.



Use case 4: Detect bloated tables

1. pgwatch is running with table-level statistics enabled.
2. Examine dashboards:
 - Tables Top → Identify tables with high dead tuples.
 - Table Details → Inspect individual table statistics.
 - Checkpointer / Bgwriter / Block IO Stats → Assess I/O impact.
3. Look for high dead tuples and table size growth.
4. Decide on action:
 - VACUUM / VACUUM FULL if autovacuum isn't sufficient.
 - REINDEX if index bloat is detected.
 - pg_squeeze for online reclaiming of space.

Use case 5: Investigate lock contention

1. pgwatch is running with lock monitoring enabled.
2. Examine dashboards:

Lock Details → Identify locked and blocking queries.

Stat Activity (Realtime) → View active queries and waiting states.

Sessions Overview → Check session and connection states.

3. Look for queries holding locks too long.
4. Resolve the issue:

Optimize queries to reduce lock duration.

Kill problematic sessions (pg_terminate_backend).

Adjust transaction isolation levels if needed.



Use case 6: Detect replication lag

1. pgwatch is running with replication monitoring enabled.

2. Examine dashboards:

Replication → Check replication slots, lag in bytes/time.

3. Look for increasing lag trends.

4. Identify causes & resolve:

Check network performance.

Tune replication parameters (wal_keep_size, max_wal_senders).

Use case 7: Analyzing Index Efficiency

1. pgwatch is running with index monitoring enabled.
2. Examine dashboards:
 - Index Overview → Identify frequently and rarely used indexes.
 - Tables Top → Detect large tables where indexing is crucial.
 - Table Details → Review index sizes and hit rates.
3. Look for indexes with very few scans (potentially unnecessary) and duplicate or redundant indexes.
4. Decide on action:
 - Drop unused indexes.
 - Rebuild bloated indexes (REINDEX).
 - Consider index type changes (e.g., B-tree vs. GIN/GiST).



Use case 8: Detect pooling issues

1. pgwatch is running with PgBouncer metrics enabled.
2. Examine dashboards:
 - PgBouncer Stats → Monitor connection pool usage and efficiency.
3. Look for:
 - High numbers of waiting clients in PgBouncer.
 - Too many active backend connections, reducing pooling efficiency.
 - Connections stuck in idle-in-transaction state, causing bloat.

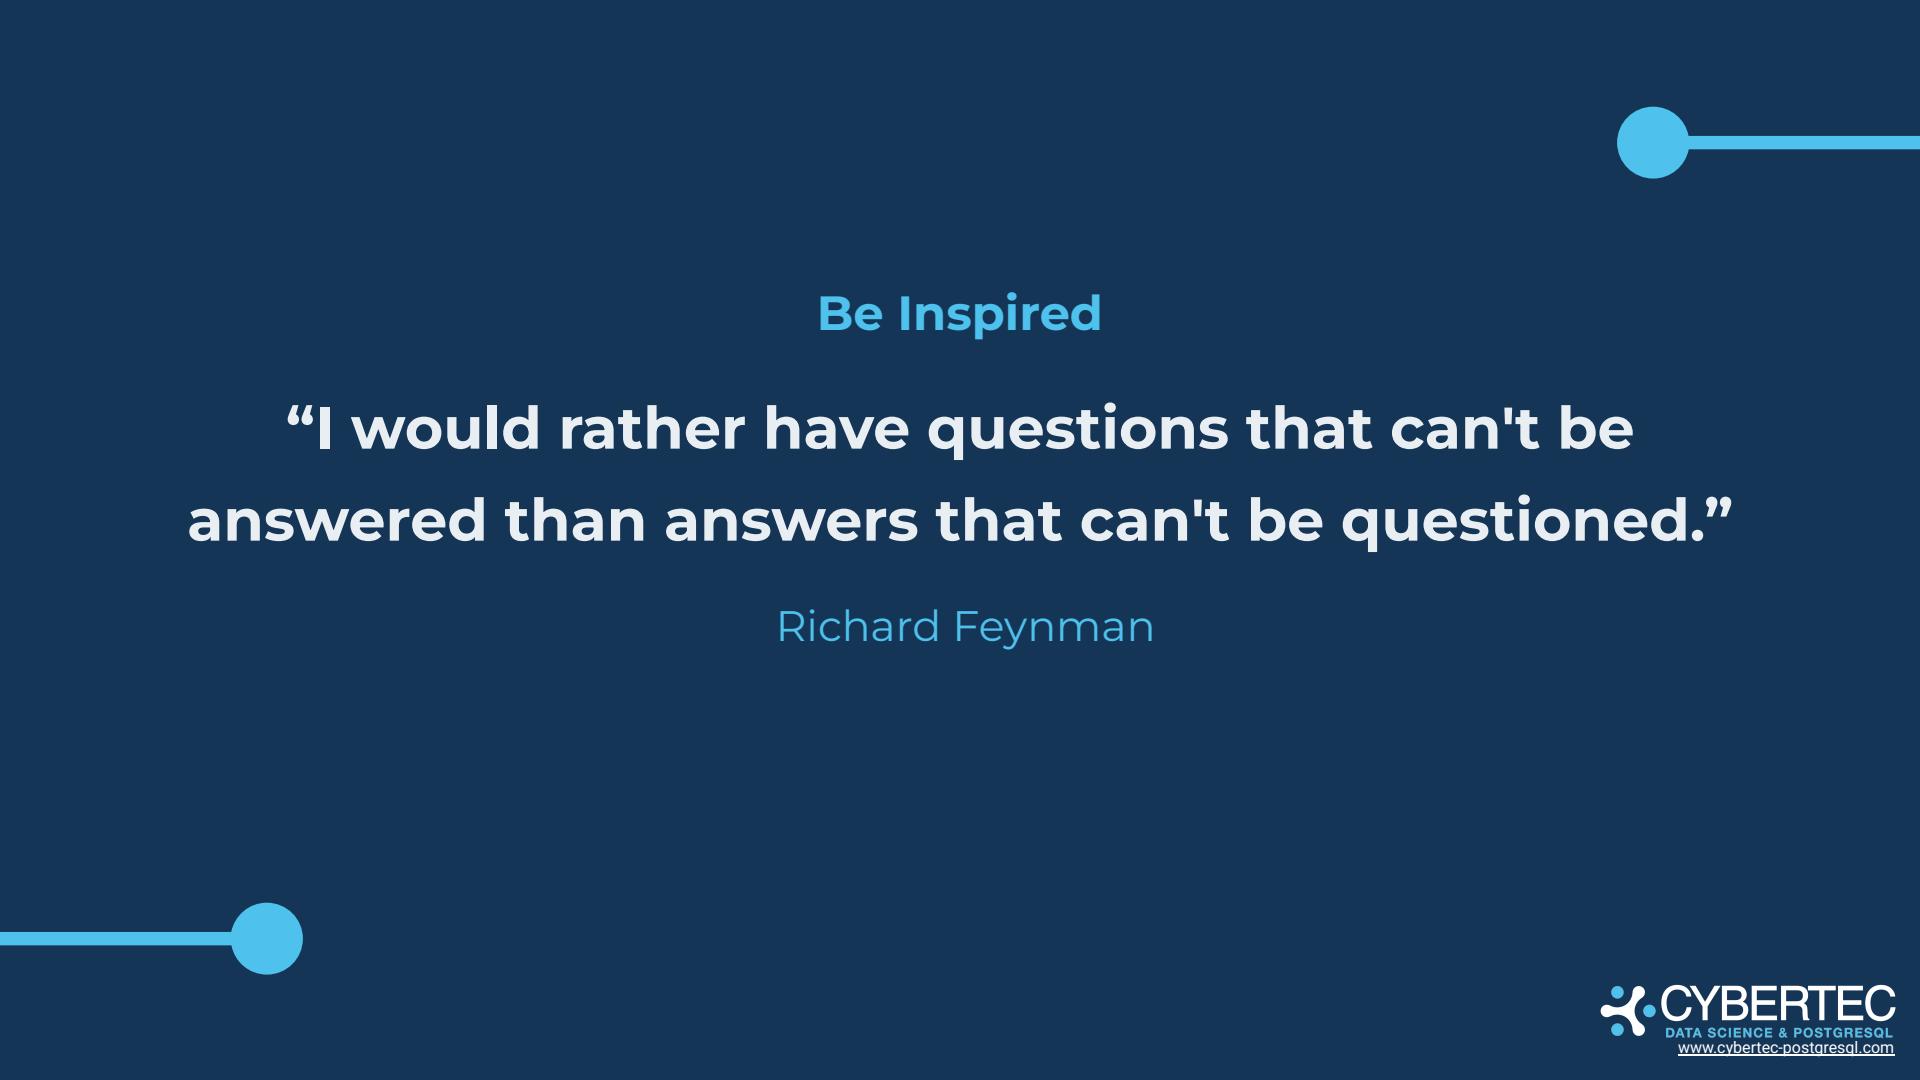
Decide on action:

- Tune PgBouncer settings (max_client_conn, default_pool_size).
- Optimize transaction vs. session pooling.
- Adjust PostgreSQL connection limits (max_connections).



Use case 9: Detect changes in settings

1. pgwatch is running with configuration tracking enabled.
2. Examine dashboards:
 - Change Events → Look for recent modifications in PostgreSQL settings.
 - DB Overview time lag → Compare overall performance before/after.
 - System Stats → Identify CPU/memory/disk usage variations.
3. Look for:
 - Critical changes (work_mem, shared_buffers, max_connections, etc.).
 - Unexpected changes (e.g., autovacuum disabled, logging settings altered).
 - Correlation between system performance and configuration changes.
4. Decide on action:
 - Verify when and who changed the settings and restore previous values
 - Implement change tracking (e.g., version control for postgresql.conf).



Be Inspired

**“I would rather have questions that can't be
answered than answers that can't be questioned.”**

Richard Feynman

DON'T BE A STRANGER



PERSONAL PAGE

pashagolub.github.io



PERSONAL GITHUB

www.github.com/pashagolub



CYBERTEC BLOG

www.cybertec-postgresql.com/en/blog/



CYBERTEC GITHUB

www.github.com/cybertec-postgresql

#StandWithUkraine

