



# Branching in PostgreSQL

Noémi Ványi (@kvch)

Prague PostgreSQL Meetup: November Edition

# Content review

**01** Branch definition

**02** Isolation options

**03** Schema level isolation

**04** Shared cluster

**05** Takeaways, questions

# A branch is an isolated database environment

to test schema changes safely

# Database environment

---

- Cluster level objects
  - Roles
  - Extensions
- Database level objects
  - Schemas
- Schema level objects
  - Tables
  - Sequences
  - ...

# Isolation options

How can we isolate branches?

## OPTION 1

**SCHEMA LEVEL**

A branch is a collection of database schemas. All users live in the same PostgreSQL database on the same PostgreSQL cluster.

## OPTION 2

**DATABASE LEVEL**

A branch is a PostgreSQL database. Multiple users can live on the same PostgreSQL cluster.

## OPTION 3

**CLUSTER LEVEL**

A branch is a PostgreSQL cluster.

infrastructure and ops effort

development effort

# Schema level isolation: Lightweight branches



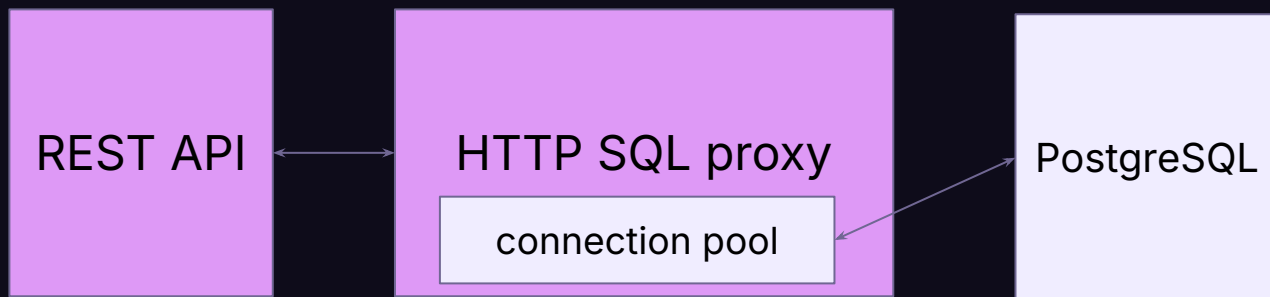
## Why opt for schema level isolation?

- Existing user branches were already in their own schemas
- No need to change our architecture significantly
- We only wanted to support update, insert, delete, truncate statements (or so we thought)
- Everything else was done by Xata REST API

# Shared clusters

Our shared PostgreSQL clusters

# SQL over HTTP



## SQL proxy over HTTP

- In the connection pool we used a common role
- When a user connected, we set the role their branch owner role
- We had to control what statements users can run
  - SET and RESET ROLE is strictly forbidden
- We forbid several functions that could leak the internals
- Limited experience because we did not support transactions or anything session level

## Tasks of the SQL proxy


- User submits their SQL statement using the REST endpoint
- SQL proxy parses the statement and decides if it is allowed
- If the statement has a syntax error or is forbidden, proxy returns an error
- If the statement is allowed, the proxy runs the statement in PostgreSQL
- It parses the result and returns it to the user in JSON or in array format

**It's all fun and games until  
someone wants to support  
wire protocol access**

## Wire protocol access use cases

- psql support is required
  - Create and manage tables, schemas
  - Export and import databases (pg\_dump and pg\_restore)
  - Stop long running queries
- ORM support
  - Several users interact with Xata using Drizzle

# SQL wire proxy





## Hiding database objects from other users

- When you list the schemas in a database, you can see the list of all schemas including the ones that you do not own or have no access to
- PostgreSQL lets you lookup object names based on OID
- With functions `regclassout`, `regclass` and `to_regclass` you can inspect objects in the database

## Shadow catalog

---

- New schema named xata\_catalog (invisible to users)
- Store all overwritten pg\_catalog functions we need
- When SQL proxy detects a call to a catalog function, it checks if it has to be rewritten to our catalog function. If yes, the function in our shadow catalog is called.
- Otherwise the function is passed through as is

## Enforce usage quotas on shared clusters

- Limit the number of concurrent connections to PostgreSQL in the gateway (same as HTTP proxy)
- With long running sessions, we decided to limit the session time length
- SQL parser found session timeout settings and if a user tried to set higher than the allowed time, we returned an error and asked the user to decrease the timeout

## Multiple schema support in a branch

Xata branch = PostgreSQL schema

Schema name in PostgreSQL is branch ID

Example: bb\_whateverid

## Multiple schema support in a branch

Xata branch = set of PostgreSQL schemas

Schema name in PostgreSQL is branch ID + schema name suffix

Example: bb\_whateverid = public schema

bb\_whateverid\_dev = dev schema

## Multiple schema support in a branch

---

- Whenever a user interacts with their database objects in a schema, we need to rewrite the schema name from the virtual (user defined) to the physical schema name
- We rewrite all statements where a schema name is present

## Multiple schema support in a branch

- We had to rewrite the outgoing names as well
- When `pg_dump` exports the database, we must return virtual schema names
- We overwrote two PostgreSQL functions `_public_nspname` and `_public_nspname_sql_identifier`
- If the schema name started with the branch prefix stored in `xata.physical_schema_prefix` or with `bb_` prefix we stripped the branch ID or returned `public`
- Rewrite output of more functions: `format_type`, `pg_get_expr`, etc.

## Multiple schema support in a branch

- When user sets `search_path`, set virtual search path in `xata.search_path` and store physical search path in `search_path`
- Read virtual search path from `xata.search_path` path and return it
- Override several function to return correct schema names like `current_schema` or `current_schemas`



## Statements we could not support

---

- DO: limited support due to schema name rewriting
- Functions: limited support, requires plpgsql parsing and rewriting
- RESET ALL
- DISCARD ALL

# Takeaways

